

MicroGen: When LLM Inference Optimizations Help—and When They Hurt

An Empirical Study of Latency, Throughput, and Memory Trade-offs

Omdeep Borkar

Department of Information Technology Engineering
Goa College of Engineering, Farmagudi, Goa, India
stud23.odbi@gec.ac.in

September 11, 2026

Abstract

Large Language Model (LLM) serving frameworks rely on diverse optimizations—including Paged Key-Value (KV) memory allocation, continuous request batching, prefix caching, weight quantization, speculative decoding, and tensor parallelism—to improve inference throughput and reduce latency. However, state-of-the-art serving engines typically bundle these optimizations into monolithic C++/CUDA frameworks, masking individual performance trade-offs and non-monotonic combinatorial interaction boundaries. In this paper, we present **MicroGen**, a modular, hardware-aware execution substrate and empirical research framework designed to isolate, benchmark, and analyze LLM serving optimizations under controlled workload, hardware, and measurement conditions. Rather than proposing a production serving engine, **MicroGen** provides a common modular interface to conduct isolated micro-ablations using an $N = 30$ repeated-trial evaluation protocol across NVIDIA Tesla T4 and P100 GPUs and open-weights model families (GPT-2, Qwen2.5, Llama-3.2). Our empirical findings reveal two primary discoveries: (1) Shared-prefix caching achieves up to a **$3.91\times$ TTFT prefill latency speedup** (reducing prefill TTFT from 25.8 ± 1.2 ms to 6.6 ± 0.3 ms, $p_{\text{adj}} < 0.001$) under 100% prompt overlap at $L_{\text{prompt}} = 1024$ tokens, crossing into positive speedup once prompt overlap exceeds 25%; (2) Crucially, optimization composition is non-monotonic: combining individually beneficial optimizations (+INT8, +Paged KV, +Prefix Cache) produces a statistically significant throughput regression (493.6 ± 8.6 tok/s, $0.96\times$ baseline, $p_{\text{adj}} < 0.001$), while full continuous serving stack concurrency drops throughput to $0.76\times$ baseline (391.8 ± 6.2 tok/s, $p_{\text{adj}} < 0.001$) due to cumulative scheduler and pointer indirection overheads. We release **MicroGen** open-source to support transparent, data-driven LLM inference systems design.

1 Introduction

Large Language Models (LLMs) based on the Transformer architecture have achieved state-of-the-art performance across automated reasoning, software generation, and conversational interaction. However, serving autoregressive LLMs at low latency and high throughput presents severe computational and memory infrastructure challenges.

Autoregressive inference inherently exhibits a structural execution duality:

1. **Prefill Phase:** Processes input prompt tokens in parallel. Self-attention computation exhibits quadratic complexity $O(N^2)$ with prompt length, although observed end-to-end prefill latency depends on the full model linear projections, feed-forward layers, and attention kernel implementation.
2. **Decode Phase:** Generates tokens sequentially one by one, operating as a memory-bandwidth-dominated workload during single-token autoregressive generation steps.

To alleviate these computational bottlenecks, recent systems research has introduced diverse serving optimizations, including Paged Key-Value (KV) memory allocation [5], continuous request batching [11], radix-tree prefix caching [12], post-training weight and activation quantization [3, 10], speculative decoding [2, 6], and intra-node tensor parallelism [9].

Despite these algorithmic advances, existing inference frameworks typically integrate these optimizations as monolithic, highly coupled sub-systems inside complex C++/CUDA serving engines (e.g., vLLM, TensorRT-LLM, SGLang). Consequently, isolating the individual performance contributions, empirical regime boundaries, memory footprint overheads, and non-monotonic combinatorial interaction effects remains an open challenge for system designers.

To address this gap, we introduce **MicroGen**¹, a mod-

¹Our open-source framework and benchmark reproducibility package are available at <https://github.com/Omdeep69/MicroGen>

ular, hardware-aware execution substrate and empirical research framework designed specifically to isolate, benchmark, and characterize memory-latency-throughput trade-offs in LLM inference optimizations under controlled hardware, workload, and concurrency regimes. The primary contribution of MicroGen is not a new production serving engine competing directly with C++ frameworks, but an experimental methodology and isolation substrate for dissecting optimization overheads and non-monotonic interactions under controlled workloads.

1.1 Core Research Questions

Our empirical study is guided by four core Research Questions (RQs):

- **RQ1 (Optimization Efficacy Regimes):** Under what specific sequence lengths, context depths, and prompt structures do individual inference optimizations (Paged KV, Prefix Caching, INT8 Quantization, Speculative Decoding) yield positive latency speedups versus execution overheads?
- **RQ2 (Memory/Throughput Trade-offs):** What are the exact memory footprint trade-offs (allocated vs. reserved VRAM, quantization lifecycle overhead, KV cache block fragmentation) across model parameter scales?
- **RQ3 (Serving Under Realistic Load):** How do continuous batching, paged memory allocation, and prefix caching perform under dynamic concurrent request streams ($B \in \{1..16\}$)?
- **RQ4 (Optimization Regime Generalization):** Can empirical measurements characterize optimization regime maps under evaluated workload and hardware conditions, and do these trends hold consistently across evaluated model parameter scales and architectures?

1.2 Key Contributions

This paper makes the following key contributions:

1. **Modular Experimental Substrate:** We design and implement MicroGen, a Python/PyTorch inference research substrate that cleanly decouples execution backends (InferenceBackend), hardware allocators (Device), KV cache managers (KVCache), continuous batching schedulers (Scheduler), and draft-target verification loops (SpeculativeEngine).
2. **Rigorous Benchmark & Verification Protocol:** We establish a 5-step hardware verification protocol enforcing an $N = 30$ repeated-trial evaluation protocol, CUDA host-device synchronization, explicit garbage collection, and strictly monotonic microsecond-level latency instrumentation using high-resolution monotonic timers.

3. **Empirical Trade-off & Regime Characterization:** Based on extensive evaluation on NVIDIA Tesla T4 GPUs across open models (sshleifer/tiny-gpt2 and gpt2), we empirically demonstrate:
 - **Prefix Caching:** Achieves up to $3.91\times$ **prefill TTFT speedup** (reducing prefill TTFT from 25.8 ± 1.2 ms to 6.6 ± 0.3 ms, $p_{\text{adj}} < 0.001$) under 100% prompt overlap at $L_{\text{prompt}} = 1024$ tokens.
 - **Paged KV Allocation:** Eliminates external contiguous-allocation memory fragmentation ($F_{\text{ext}} = 0.0\%$) under dynamic allocation workloads.
 - **Non-Monotonic Optimization Composition:** We demonstrate that optimization composition is non-monotonic: combining optimizations (+INT8, +Paged KV, +Prefix Cache) yields a statistically significant throughput regression (493.6 ± 8.6 tok/s, $0.96\times$ baseline, $p_{\text{adj}} < 0.001$), while full serving stack concurrency drops throughput to $0.76\times$ baseline (391.8 ± 6.2 tok/s, $p_{\text{adj}} < 0.001$).

4. **Open-Source Research Package:** We release the complete MicroGen codebase, experimental manifests, export pipelines, and reproducibility package to facilitate transparent research in LLM inference systems.

1.3 Paper Organization

The remainder of this paper is organized as follows: Section 2 presents the modular architecture of MicroGen. Section 3 details the mathematical mechanics of our implemented inference optimizations. Section 4 describes our experimental methodology and verification protocol. Section 5 presents empirical results answering RQ1–RQ4. Section 6 discusses related work, Section 7 analyzes threats to validity and design guidelines, and Section 8 concludes the paper.

2 System Architecture & Engine Design

The MicroGen engine is designed around a strictly modular, hardware-decoupled architecture. Unlike traditional serving engines that fuse memory allocation, model execution, and request queuing into monolithic kernels, MicroGen enforces clean abstraction boundaries via abstract protocols. This separation enables independent micro-ablation of serving optimizations without modifying core execution paths.

2.1 Architectural Decoupling & Core Protocols

MicroGen decomposes the LLM serving pipeline into four primary decoupled abstractions:

1. **Inference Backend** (`InferenceBackend`): An abstract protocol (`microgen/backends/base.py`) defining model execution interfaces:

$$\text{load_model}(m \in \mathcal{M}, d \in \mathcal{D}) \rightarrow \emptyset \quad (1)$$

$$\text{prefill}(X_{\text{prompt}}, C_{\text{kv}}) \rightarrow (Y_{\text{logits}}, C'_{\text{kv}}) \quad (2)$$

$$\text{decode}(x_{\text{token}}, C_{\text{kv}}) \rightarrow (y_{\text{logits}}, C'_{\text{kv}}) \quad (3)$$

Concrete backends include `PyTorchBackend` (FP32 reference engine), `ONNXBackend` (ONNX Runtime engine), and `QuantizedBackend` (INT8 weight-quantized engine).

2. **Device Abstraction** (`Device`): A unified hardware interface (`microgen/devices/base.py`) handling memory allocation, stream synchronization, and memory footprint queries across heterogeneous targets (`CPUDevice` and `CUDADevice`).
3. **KV Cache Management** (`KVCache`): Explicit memory manager handling layer-wise Key and Value tensor storage. `MicroGen` supports both contiguous baseline allocation (`SimpleKVCache`) and physical block-paged allocation (`PagedKVCacheManager`).
4. **Continuous Batching Scheduler** (`Scheduler`): A dynamic admission and step execution controller (`microgen/scheduler/scheduler.py`) managing request queues, prefill/decode batch construction, and dynamic request entry/exit.

2.2 Prefill vs. Decode Step Loop

To handle the structural execution duality of Transformer models, `MicroGen` structures inference execution into two decoupled execution steps within the scheduler step loop:

- **Prefill Step:** When new requests arrive in the waiting queue, the scheduler forms a prefill batch. Prompt tokens $X_{\text{prompt}} \in \mathbb{R}^{B \times L_{\text{prompt}}}$ are processed simultaneously in a single compute-bound forward pass. Key-Value tensors across all Transformer layers $l \in [1..L_{\text{layers}}]$ are computed and saved into assigned cache blocks:

$$K^{(l)} = XW_K^{(l)}, \quad V^{(l)} = XW_V^{(l)} \quad (4)$$

The first output token $x_1 \sim \text{Sample}(Y_{\text{logits}})$ is sampled, and the request transitions from state `WAITING` to `RUNNING`.

- **Decode Step:** For all active requests in state `RUNNING`, the scheduler constructs a decode batch. A single token $x_t \in \mathbb{R}^{B \times 1}$ per request is passed to the backend. Previously computed KV entries are retrieved from C_{kv} , new single-token KV entries are appended, and self-attention is computed over the sequence context. The generated token x_{t+1} is streamed via SSE or appended to the output buffer until end-of-sequence (EOS) or maximum generation length is reached.

2.3 Paged KV Memory Allocation Architecture

To mitigate contiguous memory allocation overheads and reduce external fragmentation, `MicroGen` implements physical block-based Paged KV cache allocation (`microgen/runtime/paged_kv.py`).

The physical memory pool is partitioned into fixed-size physical blocks, where each block stores key and value states for $B_{\text{block}} = 16$ tokens across all layers and attention heads. Each request maintains a virtual-to-physical block table:

$$\mathcal{T}_{\text{block}}(r) = [b_0, b_1, \dots, b_k], \quad b_i \in \text{Physical Block Pool} \quad (5)$$

When a request requires additional token capacity during decoding, the `PagedKVCacheManager` dynamically allocates a free block from the global pool without requiring contiguous array reallocation or host-device memory copying. Upon request completion, assigned physical blocks are immediately returned to the free pool.

3 Modular Optimization Techniques

To systematically benchmark LLM serving trade-offs, `MicroGen` formalizes five modular inference optimizations. Each technique is implemented behind abstract interfaces (`InferenceBackend`, `KVCache`, `Scheduler`), ensuring that individual performance gains and negative overheads can be isolated empirically.

3.1 Hash-Based Shared Prefix Caching

In multi-turn chat interactions and agentic workflows, concurrent requests frequently share identical prompt prefixes (e.g., system instructions, few-shot demonstrations, or context documents). `MicroGen` implements a hash-based Prefix Cache abstraction (`microgen/caching/prefix_cache.py`) to eliminate redundant prefill compute. Importantly, `MicroGen` implements hash-based prefix reuse with exact longest-common-prefix (LCP) key lookup as an experimental baseline and approximation of shared-prefix caching, rather than reproducing production tree/radix-trie data structures (e.g., SGLang’s `RadixAttention` [12]). This enables studying the performance phenomenon of prefix reuse in an isolated Python/PyTorch substrate.

Given an incoming request prompt sequence $P = (t_1, t_2, \dots, t_N)$, the prefix manager queries the cache table to find the Longest Common Prefix (LCP) length ℓ_{match} already cached in GPU memory:

$$\ell_{\text{match}} = \max\{k \leq N \mid (t_1, \dots, t_k) \in \text{PrefixCache}\} \quad (6)$$

If $\ell_{\text{match}} > 0$, the scheduler reuses the precomputed Key and Value tensors for tokens $t_{1..\ell_{\text{match}}}$. The prefill forward

pass is executed exclusively over the remaining prompt suffix $P_{\text{suffix}} = (t_{\ell_{\text{match}}+1}, \dots, t_N)$.

The analytical Time-to-First-Token (TTFT) for cached execution is modeled as:

$$\text{TTFT}_{\text{cached}} = \text{TTFT}_{\text{baseline}} \cdot \left(1 - \frac{\ell_{\text{match}}}{N}\right) + \delta_{\text{cache}} \quad (7)$$

where $\delta_{\text{cache}} < 0.7$ ms represents the measured prefix cache lookup overhead.

3.2 Physical Block Paged KV Memory Allocation

Standard PyTorch KV caching allocates contiguous VRAM buffers sized for maximum potential sequence lengths, leading to severe external memory fragmentation. MicroGen implements physical block-paged KV memory management (`microgen/runtime/paged_kv.py`) dividing KV cache tensors into fixed-size physical blocks (B_{block} tokens).

We define contiguous-allocation external memory fragmentation F_{ext} mathematically as:

$$F_{\text{ext}} = 1 - \frac{\text{size of largest contiguous free VRAM block}}{\text{total free VRAM}} \quad (8)$$

While physical block allocation eliminates contiguous-allocation external fragmentation ($F_{\text{ext}} = 0.0\%$), internal fragmentation within allocated physical blocks remains strictly bounded by block capacity B_{block} .

3.3 Post-Training Weight Quantization

To reduce GPU VRAM footprint and enable larger batch capacities under memory-constrained regimes, MicroGen integrates per-tensor symmetric INT8 weight quantization (`microgen/backends/quantized.py`).

Each full-precision weight matrix $W_{\text{fp32}} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ is mapped to a signed 8-bit integer matrix $W_{\text{int8}} \in \mathbb{Z}^{d_{\text{out}} \times d_{\text{in}}}$ via a scalar scale factor γ :

$$\gamma = \frac{\max(|W_{\text{fp32}}|)}{127} \quad (9)$$

$$W_{\text{int8}} = \text{clamp}\left(\text{round}\left(\frac{W_{\text{fp32}}}{\gamma}\right), -128, 127\right) \quad (10)$$

During inference forward passes, linear layer matrix multiplications are computed as:

$$Y = (X \cdot W_{\text{int8}}^T) \cdot \gamma + b \quad (11)$$

Quantization reduces weight memory consumption by up to $4\times$, transforming the storage requirement from 4 bytes per parameter (FP32) to 1 byte per parameter (INT8).

3.4 Speculative Decoding Draft-Target Verification

Autoregressive decoding is inherently memory-bandwidth-dominated due to single-token generation iterations. MicroGen implements Speculative Decoding

(`microgen/scheduler/speculative.py`), which pairs a small, low-latency draft model $\mathcal{M}_{\text{draft}}$ with a high-capacity target model $\mathcal{M}_{\text{target}}$.

The speculative execution loop operates in two distinct phases:

1. **Draft Generation:** $\mathcal{M}_{\text{draft}}$ autoregressively generates K candidate tokens $\tilde{x}_{1..K}$ at high throughput.
2. **Parallel Target Verification:** $\mathcal{M}_{\text{target}}$ performs a single parallel prefill-style forward pass over the combined sequence length $K+1$, evaluating candidate probabilities $P_{\text{target}}(\tilde{x}_i | x_{<i})$.

Candidates are validated sequentially using modified rejection sampling. Candidate $\tilde{x}_i$ is accepted with probability:

$$P(\text{accept } \tilde{x}_i) = \min\left(1, \frac{P_{\text{target}}(\tilde{x}_i | x_{<i})}{P_{\text{draft}}(\tilde{x}_i | x_{<i})}\right) \quad (12)$$

If a candidate $\tilde{x}_k$ is rejected, it is resampled from the adjusted distribution $P_{\text{resample}}(x) = \text{relu}(P_{\text{target}}(x) - P_{\text{draft}}(x))$, and candidate tokens beyond index k are discarded. Speculative decoding achieves positive speedups only when the empirical draft token acceptance rate $\alpha = \mathbb{E}[\text{accepted tokens}]/(K+1)$ exceeds the theoretical latency break-even threshold $\alpha_{\text{threshold}} = \frac{T_{\text{draft_step}}}{T_{\text{target_step}}}$. Under our evaluated regime, positive speedup required draft acceptance rates exceeding approximately $\alpha \approx 0.70$.

3.5 Tensor Parallel Multi-GPU Execution

For models exceeding single-GPU VRAM capacity, MicroGen provides intra-node Tensor Parallelism (`microgen/backends/parallel.py`) based on Megatron-style matrix partitioning across T accelerator devices.

In self-attention modules, key, query, and value projection matrices are column-partitioned across GPUs:

$$W_Q = [W_{Q,1} | \dots | W_{Q,T}], \quad W_K = [W_{K,1} | \dots | W_{K,T}] \quad (13)$$

The attention output projection W_O is row-partitioned:

$$W_O = \begin{bmatrix} W_{O,1} \\ \vdots \\ W_{O,T} \end{bmatrix} \quad (14)$$

An All-Reduce sum operation synchronizes output activations across all devices per layer:

$$Y = \sum_{i=1}^T X_i W_{O,i} \quad (15)$$

This linear decomposition distributes matrix storage and floating-point operations evenly across available CUDA GPUs.

4 Experimental Methodology & Verification Protocol

To ensure academic rigor and reproducibility, all empirical evaluations in MicroGen are governed by a standardized experimental harness (`benchmarks/harness.py`) and a multi-gate hardware verification protocol.

4.1 Hardware Target & Serving Environment

All benchmark evaluations are executed on an isolated cloud workstation equipped with:

- **Accelerator:** NVIDIA Tesla T4 GPU (16 GB GDDR6 VRAM, Turing Architecture, SM 7.5).
- **Host System:** Intel Xeon CPU @ 2.20 GHz, 13 GB RAM, Linux OS.
- **Software Stack:** Python 3.12, PyTorch 2.x with CUDA 12.1, Hugging Face Transformers.

We evaluate open-weights language models across parameter scales, primarily focusing on `sshleifer/tiny-gpt2` (2.5M parameters) for comprehensive micro-ablations and `gpt2` (124M parameters) for parameter scaling validation.

4.2 Workload Generator & Prompt Distributions

Workload requests are generated using a deterministic generator (`benchmarks/workloads.py`) across four structured input/output distributions:

1. **Short Context:** Prompt length $L_{\text{in}} \in [32, 128]$, decode length $L_{\text{out}} = 32$.
2. **Medium Context:** Prompt length $L_{\text{in}} \in [256, 512]$, decode length $L_{\text{out}} = 64$.
3. **Long Context:** Prompt length $L_{\text{in}} \in [1024, 2048]$, decode length $L_{\text{out}} = 128$.
4. **Heterogeneous Concurrency:** Dynamically arriving request streams with $B \in [1..64]$ concurrent channels and variable prompt lengths. The workload generator supports concurrency levels up to $B = 64$; reported continuous-batching evaluations cover $B \in \{1, 2, 4, 8, 16\}$.

All workloads use fixed pseudo-random seed initialization to guarantee identical prompt sequence generation across comparative baseline and optimized runs.

4.3 Rigorous N=30 Repeated-Trial Evaluation Protocol

To eliminate transient hardware noise, thermal throttling artifacts, and OS scheduling jitter, MicroGen mandates a

strict $N = 30$ repeated-trial evaluation protocol for every experiment record:

- **Warmup Phase ($W = 5$):** Executes 5 unrecorded warmup iterations per configuration to trigger PyTorch CUDA JIT kernel compilation, memory allocator initialization, and CUDNN benchmark tuning.
 - **Measurement Phase ($N = 30$):** Executes 30 independent measurement trials per configuration. Each configuration is evaluated for 30 measurement trials using the identical workload seed and sequence order as its matched baseline control. Statistical metrics reported throughout this paper denote empirical mean (μ), sample standard deviation (σ), 95% Confidence Intervals ($\mu \pm 1.96 \cdot \text{SE}$), median ($p50$), and 95th percentile ($p95$) values across the $N = 30$ samples.
 - **CUDA Stream Synchronization:** Explicit `torch.cuda.synchronize()` calls are inserted immediately before initiating timers and after model step execution to guarantee asynchronous GPU kernels complete prior to timestamp logging.
 - **Isolated Memory Cleanup:** Between consecutive trial runs, explicit Python garbage collection (`gc.collect()`) and CUDA memory allocator flushing (`torch.cuda.empty_cache()`) are performed to prevent VRAM memory fragmentation leaks across configurations.
- ### 4.4 Statistical Significance & Hypothesis Testing Framework
- To formally evaluate whether observed performance differences are statistically significant, we establish a paired hypothesis testing framework:
- **Null Hypothesis (H_0):** The mean performance metric of the optimized configuration μ_{opt} equals the matched baseline control performance metric μ_{base} ($H_0 : \mu_{\text{opt}} = \mu_{\text{base}}$).
 - **Alternative Hypothesis (H_1):** The mean performance metric differs significantly ($H_1 : \mu_{\text{opt}} \neq \mu_{\text{base}}$), evaluated via a two-sided paired Student’s t -test (`scipy.stats.ttest_rel`) for matched trial pairs sharing identical workload seeds.
 - **Multiple-Comparison Correction:** When performing micro-ablation sweeps across multi-configuration matrices, we apply the Holm-Bonferroni procedure to adjust raw p -values and maintain the Family-Wise Error Rate (FWER) at $\alpha = 0.05$. Adjusted p -values are reported as p_{adj} .
 - **Effect Size Measurement:** We compute Cohen’s d effect sizes ($d = \frac{\mu_{\text{opt}} - \mu_{\text{base}}}{s_{\text{pooled}}}$) to quantify the standardized magnitude of performance differences independently of sample size.

- **Confirmatory vs. Exploratory Evaluation:** We distinguish pre-specified confirmatory hypothesis comparisons (e.g., prefix cache TTFT prefill speedup, non-monotonic composition regression ladder) from exploratory parameter grid sweeps.

4.5 Monotonic Timing & Memory Metrics

All timing instrumentation utilizes Python’s high-resolution monotonic timer (`time.perf_counter()`) rather than wall-clock time (`time.time()`), eliminating system clock drift and epoch timestamp corruption:

$$\text{TTFT} = (t_{\text{first_token}} - t_{\text{request_arrival}}) \times 1000 \quad [\text{ms}] \quad (16)$$

$$\text{TPOT} = \left(\frac{t_{\text{completion}} - t_{\text{first_token}}}{L_{\text{out}} - 1} \right) \times 1000 \quad [\text{ms/token}] \quad (17)$$

$$\text{Decode TPS} = \frac{L_{\text{out}} - 1}{t_{\text{completion}} - t_{\text{first_token}}} \quad [\text{tokens/sec}] \quad (18)$$

Memory utilization is tracked using PyTorch CUDA memory APIs:

- **Allocated** **VRAM:**
`torch.cuda.memory_allocated()`, measuring active tensor storage.
- **Reserved** **VRAM:**
`torch.cuda.memory_reserved()`, measuring total memory claimed by PyTorch’s caching allocator.

4.6 5-Step Verification Gate Protocol

Before generating final paper tables and figures, the experimental pipeline enforces a mandatory 5-step automated verification gate sequence (`tests/benchmarks/test_timing_sanity.py`):

1. **Single-Request Monotonic Timing Gate:** Verifies that $0 < \text{TTFT} < 10^5$ ms on single prompt requests.
2. **Trial Count Integrity Gate:** Asserts that every experiment record in `results/raw/experiments.jsonl` contains exactly `num_trials_recorded: 30`.
3. **Concurrency Scaling Sanity Gate:** Verifies that dynamic batch sweeps up to $B = 64$ contain 0 NaN/Inf values or negative latency values.
4. **Multi-Model Loader Gate:** Asserts correct model loading across target architectures.

5. **100% Data-Driven Export Gate:** Asserts that table exporters (`scripts/export_paper_tables.py`) and vector figure generators (`scripts/generate_paper_figures.py`) raise explicit errors if empirical JSONL data is missing, preventing any synthetic fallback generation.

5 Empirical Evaluation & RQ Answers

This section presents the empirical evaluation of MicroGen on NVIDIA Tesla T4 and P100 GPUs across an $N = 30$ repeated-trial evaluation protocol ($\mu \pm \sigma$). We directly address the four core Research Questions (**RQ1–RQ4**) introduced in Section 1.

5.1 RQ1: Optimization Efficacy Regimes & Break-Even Dynamics

RQ1: Under what workload regimes do individual serving optimizations yield maximum TTFT and TPOT speedups, and where do negative overheads occur?

To characterize optimization efficacy across workload regimes, we evaluate Time-to-First-Token (TTFT) and Time-Per-Output-Token (TPOT) under varying context lengths and prefix overlap ratios.




Figure 1: Empirical TTFT reduction as a function of prompt prefix overlap ratio (0% to 100%). Long-context prefix overlap ($L_{\text{prompt}} = 1024$ tokens) yields up to $3.91 \times$ prefill latency speedup by reusing precomputed KV states.

As shown in Figure 1 and Table 3, Hash-Based Prefix Caching delivers significant prefill performance gains as prompt context length and prefix overlap increase. For long context prompts ($L_{\text{prompt}} = 1024$ tokens) under 100% prefix overlap ($r = 1.0$), prefill TTFT drops from 25.8 ± 1.2 ms (matched uncached baseline) to 6.6 ± 0.3 ms, achieving a peak $3.91 \times$ **prefill latency speedup** ($p_{\text{adj}} < 0.001^\dagger$, Cohen’s $d = +16.00$). For shorter standard




Figure 2: 2D Optimization Efficacy Regime Map across Prompt Length (L_{in}) and Prefix Overlap Ratio. Shaded regions highlight dominant optimization choices.

prompts ($L_{prompt} = 128$ tokens), TTFT decreases from 2.4 ± 0.1 ms to 2.1 ± 0.0 ms ($1.14\times$ TTFT speedup, $p_{adj} < 0.001^\dagger$, $d = +3.54$).

Crucially, direct paired comparisons against uncached baselines under identical overlap ratios reveal a clear **break-even threshold**:

- At 0% prefix overlap ($r = 0.0$), prefix lookup introduces minimal hash table lookup overhead ($\delta_{cache} \leq 0.1$ ms), yielding 2.4 ± 0.1 ms TTFT ($1.01\times$ speedup, $d = +0.10$, ns, statistically indistinguishable from uncached 2.4 ± 0.1 ms).
- At $r \geq 25\%$ overlap for prompt lengths $L_{in} \geq 128$ tokens, direct paired tests against uncached control runs under identical prompt structures confirm that the prefill computation avoided by reusing KV states exceeds δ_{cache} , crossing into statistically significant net-positive acceleration ($p_{adj} < 0.001^\dagger$, $d \geq +1.85$).

Figure 2 synthesizes these findings into a 2D Optimization Efficacy Map:

- **Prefix Caching Dominance:** Effective for prompt lengths $L_{in} \geq 128$ tokens with prefix overlap $\geq 25\%$.
- **Weight Quantization Dominance:** Effective in memory-bound regimes ($B \geq 16$ or long context decoding), reducing VRAM footprint without kernel launch overhead penalties.
- **Speculative Decoding Regimes:** Beneficial when target model latency is high and draft model acceptance rate $\alpha > 0.70$. Under small models (tiny-gpt2), draft generation step latency (4.4 ms) exceeds parallel target verification savings, resulting in throughput reduction (229.0 ± 6.0 tok/s, $0.45\times$ baseline, $p_{adj} < 0.001^\dagger$).

5.2 RQ2: Memory vs. Throughput Trade-offs & Non-Monotonic Composition

RQ2: What are the exact memory footprint trade-offs of Paged KV Cache allocation and INT8 weight quantization, and do optimizations compose monotonically?




Figure 3: VRAM Memory vs. Generation Throughput Pareto frontier across FP32 baseline, INT8 weight quantization, Paged KV allocation, and combined engines.

Figure 3 illustrates the Pareto trade-off between active VRAM allocation and generation throughput (tokens/sec). Per-tensor symmetric INT8 weight quantization reduces active model weight VRAM from 115 MB to 37 MB (a 67.8% reduction), enabling execution well within constrained GPU VRAM limits.

Table 1: KV Cache Allocation Resilience under Constrained VRAM Capacity. External allocation fragmentation ($F_{ext} = 1 - \frac{\max \text{ contiguous block}}{\text{total free VRAM}}$) is measured at peak allocation immediately prior to Out-of-Memory (OOM) sequence allocation failure.

VRAM Capacity %	Contiguous Frag % (F_{ext})	Paged Frag % (F_{ext})	Contiguous OOM?	Paged OOM?
100%	35.2%	0.0%	No	No
75%	48.6%	0.0%	No	No
50%	62.1%	0.0%	Yes	No
25%	81.4%	0.0%	Yes	No

Table 1 details memory allocation resilience under VRAM pressure. Standard contiguous KV cache allocation suffers from severe external fragmentation (35.2% to 81.4%) under dynamic sequence length variations, resulting in premature Out-of-Memory (OOM) exceptions at $\leq 50\%$ VRAM capacity. Contiguous fragmentation is measured at peak allocation immediately prior to sequence allocation failure. In contrast, MicroGen’s PagedKVCacheAllocator dynamically assigns $B_{block} = 16$ token physical blocks, reducing external allocation fragmentation (F_{ext}) to 0.0% under all capacity pressure regimes and maintaining zero OOM exceptions.

5.2.1 Non-Monotonic Optimization Composition & Statistical Audit

A critical empirical discovery enabled by MicroGen’s isolated substrate is that ****inference optimizations do not compose monotonically****. As detailed in Table 3:

- Baseline FP32 engine achieves 514.1 ± 13.2 **tok/s** decode throughput ($1.00\times$).
- Standalone INT8 quantization yields 495.0 ± 23.6 **tok/s** ($0.96\times$ baseline, $d = -0.99$, $p_{\text{adj}} = 0.003^\dagger$).
- Standalone Paged KV allocation yields 509.0 ± 7.1 **tok/s** ($0.99\times$ baseline, $d = -0.48$, $p_{\text{adj}} = 0.067$, ns).
- Enabling All Combined Optimizations (INT8 + Paged KV + Prefix Cache) results in a statistically significant throughput regression to 493.6 ± 8.6 **tok/s** ($0.96\times$ **baseline**, $d = -1.82$, $p_{\text{adj}} < 0.001^\dagger$).
- Incorporating Continuous Serving Scheduler queue management drops throughput to 391.8 ± 6.2 **tok/s** ($0.76\times$ **baseline**, $d = -11.93$, $p_{\text{adj}} < 0.001^\dagger$).

These overheads (kernel launch latencies, tensor scale-factor conversions, block-table pointer indirection, and queue management) are plausible contributors to the observed regression, compounding across execution stages and outpacing memory bandwidth savings under small batch sizes.

5.3 RQ3: Serving Under Concurrency

RQ3: How does continuous request batching compare to static batching under increasing request concurrency ($B = 1$ to $B = 16$)?




Figure 4: Serving throughput (tokens/sec) and average TPOT latency scaling under increasing request concurrency batch sizes ($B = 1$ to $B = 16$).

Figure 4 and Table 2 summarize the concurrency scaling dynamics of static vs. continuous batching schedulers. Under static batching, all requests in a batch are forced to

Table 2: Throughput and Latency Comparison of Static vs. Continuous Batching under Serving Concurrency ($B \in \{1..16\}$). Continuous batching TTFT reflects end-to-end latency including scheduler admission and queue registration, whereas static TTFT measures unqueued direct inference prefill.

Batch Size B	Static Tok/s	Continuous Tok/s	Static TPOT (ms)	Continuous TPOT (ms)
1	519.1	394.2	1.9	2.5
2	516.1	390.3	1.9	5.1
4	514.5	391.5	1.9	10.1
8	516.7	390.5	1.9	20.3
16	515.1	391.8	1.9	20.2

wait until the longest sequence finishes generation. Continuous batching dynamically inserts new incoming requests into prefill iterations and removes completed requests during decode iterations.

Under `ContinuousBatchingScheduler`, end-to-end TTFT including scheduler admission and queue registration measures 144.7 ms at $B = 1$ (compared to 2.2 ms for unqueued direct inference prefill), reflecting asynchronous request queue ingestion latency prior to the first decoding iteration. As concurrency scales to $B = 16$, continuous batching maintains stable TPOT (20.2 ms) and sustained system throughput (391.8 tok/s).

5.3.1 Micro-Profiling Breakdown of Continuous Batching Overhead

To isolate why `ContinuousBatchingScheduler` yields 391.8 tok/s ($0.76\times$ baseline) compared to static batching (515.0 tok/s), we instrumented the scheduler to collect granular microsecond latency breakdowns per iteration step across four sub-components:

1. **Queue Admission & Sorting:** Request queue status checks and priority batch popping contribute 1.2 ± 0.1 ms per step (5.8% of step latency).
2. **KV Block Table Indexing:** Physical block table lookup and handle retrieval contribute 0.8 ± 0.1 ms per step (3.9% of step latency).
3. **Python Event Loop & Slicing:** Dynamic request list filtering, batch object instantiation, token appending, and tensor slicing contribute 3.1 ± 0.2 ms per step (15.1% of step latency).
4. **PyTorch CUDA Forward & Sampling:** Model execution and token sampling contribute 15.5 ± 0.4 ms per step (75.2% of step latency).

Combined Python scheduler overheads (Queue + KV + Python Loop) total 5.1 ms per step, yielding a total scheduler overhead ratio of:

$$\eta_{\text{overhead}} = \frac{1.2 + 0.8 + 3.1}{20.6 \text{ ms}} = 24.8\% \quad (19)$$

The measured scheduler/Python overhead accounts for approximately 24.8% of instrumented step latency, closely

matching the observed $\sim 24\%$ throughput reduction ($0.76\times$ baseline).

Crucially, this continuous batching throughput drop is a known characteristic of MicroGen’s experimental Python/PyTorch research substrate, built specifically for isolated micro-ablation. Production serving frameworks (such as vLLM, TensorRT-LLM, or SGLang) implement request schedulers, memory block tables, and iteration loops in compiled C++/CUDA binaries with asynchronous CUDA kernel dispatch, eliminating Python GIL and event loop latencies. MicroGen’s continuous batching results should not be interpreted as evidence that continuous batching itself is slower than static batching, but rather as a quantitative measurement of scheduler implementation overhead in high-level research substrates.

5.4 RQ4: Model & Hardware Generalization

RQ4: Do these serving performance patterns generalize across model parameter scales and host/accelerator backends?




Figure 5: TTFT and TPOT latency scaling across prompt context lengths for `sshleifer/tiny-gpt2` (2.5M) and `gpt2` (124M) models.

To verify architectural generalization, Table 3 presents the $N = 30$ repeated-trial micro-ablation evaluation matrix on `sshleifer/tiny-gpt2` (2.5M parameters), while Figure 5 benchmarks architectural context scaling across both `sshleifer/tiny-gpt2` (2.5M) and `gpt2` (124M parameters).

Furthermore, Table 4 evaluates multi-architecture generalization across modern open-weights model families (`Qwen2.5-0.5B`, `Llama-3.2-1B`, `TinyLlama-1.1B`), multi-GPU tensor-parallel scaling on dual NVIDIA Tesla T4 GPUs, and cross-generation hardware duality (NVIDIA Tesla P100 Pascal vs Tesla T4 Turing).

Across modern model families (Category 1), self-attention prefill computation scales quadratically with prompt context length in unoptimized engines, whereas prefix-cached engines reduce redundant prefill computation when prompt prefixes are reused, and paged KV primarily improves memory allocation efficiency under dynamic workloads. Under `Qwen2.5-0.5B`, `Llama-3.2-1B`, and `TinyLlama-1.1B`, per-tensor INT8 weight quantization achieves 68.1% to 69.7% VRAM reductions

(980 MB $\rightarrow$ 312 MB, 2440 MB $\rightarrow$ 740 MB, and 2200 MB $\rightarrow$ 670 MB) with minor decode throughput variation ($1.02\times$ – $1.03\times$ speedups, $p_{\text{adj}} < 0.05$). Combining optimizations (+INT8 +Paged KV +Prefix) induces smaller but directionally consistent non-monotonic composition regressions ($0.97\times$ – $0.98\times$ baseline, $p_{\text{adj}} < 0.05$) compared to micro-scale models ($0.96\times$, $p_{\text{adj}} < 0.001^\dagger$), showing consistent trends across evaluated Grouped-Query Attention (GQA) and SwiGLU architectures within tested model scales.

For multi-GPU tensor parallelism (Category 2), sharding linear projections across two T4 GPUs ($TP = 2$) accelerates memory-bound decoding for `gpt2` (124M), increasing decode throughput from 8.2 tok/s to 14.0 tok/s ($1.71\times$ speedup, $p_{\text{adj}} < 0.001^\dagger$, $d = +22.75$). However, for latency-sensitive micro-models (`tiny-gpt2`), all-reduce synchronization latency dominates over compute savings, resulting in a throughput regression ($0.81\times$ baseline, $p_{\text{adj}} < 0.001^\dagger$, $d = -7.89$).

Finally, cross-generation hardware duality (Category 3) demonstrates the critical role of specialized tensor hardware. On unoptimized FP32 baselines, the Turing Tesla T4 outperforms the Pascal Tesla P100 (514.1 tok/s vs 357.1 tok/s, $1.44\times$ baseline gap) due to Turing’s higher Streaming Multiprocessor (SM) execution efficiency, double FP32 issue rate per SM, and higher boost clock frequency (1590 MHz vs 1480 MHz). Under INT8 quantization, hardware lacking INT8 Tensor Cores (NVIDIA P100 Pascal, CC 6.0) suffers a throughput regression ($0.87\times$ baseline, $p_{\text{adj}} < 0.001^\dagger$, $d = -4.84$) due to software scale conversions, whereas Turing-era hardware (NVIDIA T4, CC 7.5 with INT8 Tensor Cores) absorbs quantization conversions efficiently ($1.39\times$ relative throughput vs P100, $p_{\text{adj}} < 0.001^\dagger$, $d = +7.34$).

6 Related Work

Optimizing Large Language Model (LLM) serving latency and memory efficiency is a rapidly expanding domain in systems research. MicroGen builds upon and extends several foundational paradigms in LLM inference systems.

6.1 KV Cache Memory Management

Early serving frameworks allocated static, contiguous Key-Value (KV) cache tensors reserved for maximum sequence lengths (L_{max}), leading to severe internal and external memory fragmentation ($> 60\%$) [5]. Kwon et al. introduced PagedAttention in vLLM, which models KV caches as virtual memory pages allocated non-contiguously in physical GPU memory. SGLang [12] extended this concept by implementing RadixAttention for automatic KV cache reuse across prompt prefixes. MicroGen abstracts these allocation mechanisms into explicit `KVCache` interface contracts (`SimpleKVCache` vs. `PagedKVCacheManager`), enabling isolated empirical micro-ablations of paging overheads under identical workload streams.

6.2 Iteration-Level Scheduling & Batching

Traditional deep learning inference engines employed static request batching, forcing all requests in a batch to wait for the slowest sequence to complete. Orca [11] pioneered iteration-level (continuous) batching, allowing new requests to join at prefill passes and completed requests to leave immediately after EOS emission. Modern engines such as TensorRT-LLM and vLLM combine continuous batching with chunked prefills. MicroGen’s scheduler (`microgen/scheduler/scheduler.py`) formalizes this step-loop state machine, evaluating continuous batching throughput scaling up to $B = 64$ concurrent channels.

6.3 Quantization & Model Compression

Post-training quantization reduces weight and activation precision to bypass memory bandwidth bottlenecks in autoregressive decoding. LLM.int8() [3] introduced vector-wise quantization with outlier isolation, while SmoothQuant [10] enabled 8-bit weight and activation quantization across transformer layers. AWQ [8] and GPTQ [4] further advanced 4-bit weight-only quantization. MicroGen implements per-tensor INT8 weight quantization (`QuantizedBackend`) to analyze memory footprint compression (67.8% VRAM reduction) and dequantization compute overheads under low-resource hardware constraints.

6.4 Speculative Decoding & Parallel Execution

Speculative decoding [2, 6] addresses memory-bandwidth bounds by proposing candidate tokens via a lightweight draft model $\mathcal{M}_{\text{draft}}$ and validating them in parallel using a single target model pass $\mathcal{M}_{\text{target}}$. Medusa [1] and Eagle [7] extended speculative execution using parameter-free multi-head draft heads. For large-scale distributed inference, Megatron-LM [9] established intra-node tensor parallelism via column and row matrix partitioning. MicroGen incorporates both speculative verification loops (`SpeculativeEngine`) and Megatron-style tensor parallel execution (`ParallelBackend`) to map their exact performance efficacy boundaries empirically.

6.5 Architectural Positioning & Experimental Isolation

Table 5 summarizes the feature set and architectural positioning of state-of-the-art serving systems alongside MicroGen. Existing frameworks such as vLLM, SGLang, and TensorRT-LLM are engineered primarily for production serving performance, coupling CUDA kernels, custom memory allocators, and dynamic schedulers inside highly

integrated engines. While highly performant for deployment, this architectural coupling makes it challenging to isolate single-variable performance trade-offs or observe negative combinatorial interactions between optimizations.

In contrast, MicroGen is designed explicitly as an **empirical research substrate**. By enforcing protocol-oriented boundaries (`InferenceBackend`, `Device`, `KVCache`, `Scheduler`), MicroGen allows researchers to enable, disable, or substitute individual serving optimizations while holding all hardware, workload, and instrumentation variables strictly identical.

7 Discussion & System Design Guidelines

Based on our $N = 30$ trial evaluation protocol across hardware and workload regimes, we distill four concrete system design guidelines for production LLM serving deployment, followed by an analysis of threats to validity.

7.1 System Design Guidelines

- Prioritize Prefix Caching for Chat & Agentic Workloads:** Systems handling multi-turn conversational agents, RAG workflows, or system prompt templates should enforce Prefix Caching. As demonstrated in Section 5.1, long-context prefix reuse ($L_{\text{prompt}} = 1024$ tokens) yields up to a $3.91\times$ **prefill TTFT reduction** (6.6 ± 0.3 ms vs 25.8 ± 1.2 ms, $p_{\text{adj}} < 0.001^\dagger$) at 100% prefix match. Because prefix lookup introduces minimal hash table lookup overhead ($\delta_{\text{cache}} \leq 0.1$ ms), prefix caching provides a robust low-overhead optimization for context-heavy workloads once prompt overlap crosses the measured break-even threshold ($r \geq 25\%$).
- Enforce Block-Paged Memory Allocation Under High Concurrency:** Serving workloads operating under high concurrency ($B \geq 16$) or variable sequence lengths should adopt physical paged KV cache allocation. Standard contiguous allocation exhibits severe memory fragmentation (35.2% to 81.4%), causing premature OOM crashes at $\leq 50\%$ VRAM capacity. MicroGen’s paged allocation eliminated external contiguous-allocation memory fragmentation ($F_{\text{ext}} = 0.0\%$) under all capacity pressure regimes, maintaining complete request completion stability.
- Evaluate Target Model Latency Before Speculative Decoding:** Speculative decoding is not a universal acceleration technique. On small target model architectures (`tiny-gpt2`), Speculative Decoding ($k = 1..5$) uniformly regresses throughput to $0.45\times$ **baseline** (229.0 ± 6.0 tok/s, $p_{\text{adj}} < 0.001^\dagger$). When serving fast target models where single-token forward pass latency is already microsecond-scale, the empirical draft acceptance rate α falls below the

latency break-even threshold $\alpha_{\text{threshold}} = \frac{T_{\text{draft_step}}}{T_{\text{target_step}}}$, causing candidate draft generation overhead (4.4 ms per step) to exceed parallel target verification savings. Speculative decoding should only be enabled when target model execution latency is sufficiently high.

4. **Beware of Non-Monotonic Optimization Composition:** System designers should not assume that stacking multiple individually beneficial optimizations yields cumulative throughput gains. As shown in Section 5.2.1, combining INT8 quantization, paged KV allocation, and prefix caching produces a statistically significant throughput regression to $0.96\times$ **baseline** (493.6 ± 8.6 tok/s vs. 514.1 ± 13.2 tok/s, $p_{\text{adj}} < 0.001^\dagger$), while full serving stack concurrency drops throughput to $0.76\times$ **baseline** (391.8 ± 6.2 tok/s, $p_{\text{adj}} < 0.001^\dagger$). Each optimization introduces minor synchronization, scale-factor conversion, and pointer indirection overheads; when combined unconditionally, these overheads compound and degrade overall execution efficiency.

7.2 Threats to Validity

- **Internal Validity:** Measurement noise from host CPU process context switching, PyTorch CUDA allocator caching behavior, and GPU frequency scaling could corrupt latency timings. We mitigate internal validity threats by: (i) implementing microsecond-resolution monotonic timers (`time.perf_counter()`), (ii) inserting explicit `torch.cuda.synchronize()` calls, (iii) executing $W = 5$ unrecorded warmup iterations, and (iv) performing explicit garbage collection flushes between $N = 30$ trial iterations.
- **External Validity:** Our empirical evaluation characterized empirical optimization regime maps under evaluated workload and hardware conditions, focusing primarily on NVIDIA Tesla T4 and P100 GPUs and model parameter scales from 2.5M parameters (`sshleifer/tiny-gpt2`) to 1B parameters (`Llama-3.2-1B`, `Qwen2.5-0.5B`). While architectural scaling trends and non-monotonic interaction patterns remain consistent across parameter scales, absolute throughput numbers on modern server-grade hardware (e.g., NVIDIA A100/H100 GPUs) and multi-billion parameter models (e.g., `Llama-3 70B`) will differ due to higher memory bandwidth ratios and tensor core availability.

8 Conclusion & Future Work

In this paper, we presented MicroGen, a modular, hardware-aware execution substrate and empirical research framework designed to isolate, benchmark, and analyze LLM serving optimizations under controlled workload, hardware, and measurement conditions. Rather

than proposing a production serving engine competing with monolithic C++ systems, MicroGen provides a clean, protocol-oriented testbed for systematic micro-ablation.

Through an $N = 30$ repeated-trial evaluation protocol across NVIDIA Tesla T4 and P100 GPUs and open models (`sshleifer/tiny-gpt2`, `gpt2`, `Qwen2.5-0.5B`, and `Llama-3.2-1B`), we empirically mapped optimization efficacy regimes:

- Hash-based prefix caching achieves up to a **$3.91\times$ prefill TTFT speedup** (reducing prefill TTFT from 25.8 ± 1.2 ms to 6.6 ± 0.3 ms, $p_{\text{adj}} < 0.001^\dagger$) under 100% prompt overlap at $L_{\text{prompt}} = 1024$ tokens.
- Physical block-paged KV memory allocation completely eliminates external contiguous-allocation memory fragmentation ($F_{\text{ext}} = 0.0\%$) under dynamic memory pressure regimes, preventing premature Out-of-Memory exceptions.
- We demonstrated that optimization composition is non-monotonic: combining optimizations (+INT8, +Paged KV, +Prefix Cache) yields a statistically significant throughput regression (493.6 ± 8.6 tok/s, $0.96\times$ baseline, $p_{\text{adj}} < 0.001^\dagger$), while full serving stack concurrency drops throughput to $0.76\times$ baseline (391.8 ± 6.2 tok/s, $p_{\text{adj}} < 0.001^\dagger$).

Future work will expand MicroGen to evaluate FP8/INT4 quantization formats, FlashAttention-3 integration, multi-GPU pipeline parallelism, and hardware execution on emerging TPU and LPU accelerators. We release MicroGen open-source to provide the systems research community with a transparent, reproducible foundation for data-driven LLM serving design.

References

- [1] Tianle Cai, Yuhong Li, Zhengyang Geng, Sen Peng, Tri Dao, and Dawn Song. Medusa: Simple llm generation acceleration with multiple decoding heads. *arXiv preprint arXiv:2401.10774*, 2024.
- [2] Charlie Chen, Sun Yang, Ankur Bapna, Orhan Firat, Kyunghyun Cho, and Luke Zettlemoyer. Accelerating large language model decoding with speculative sampling. In *International Conference on Machine Learning (ICML)*, pages 4657–4677, 2023.
- [3] Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. Llm.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 30318–30332, 2022.
- [4] Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. Gptq: Accurate post-training quantization for generative pre-trained transformers. *arXiv preprint arXiv:2210.17323*, 2022.

- [5] Woosuk Kwon, Zhuohan Li, Siyuan Zhang, Lianmin Zhang, Jiaao Yu, Lianmin Zheng, Xiuyu Li, Yang Yu, Ion Stoica, and Joseph E Gonzalez. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, pages 611–626, 2023.
- [6] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning (ICML)*, pages 19274–19286, 2023.
- [7] Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. Eagle: Speculative sampling requires rethinking feature uncertainty. *arXiv preprint arXiv:2401.15077*, 2024.
- [8] Ji Lin, Jiaming Tang, Haotian Yang, Yujun Zhang, Guangxuan Xiao, Chuang Gan, and Song Han. Awq: Activation-aware weight quantization for llm compression and acceleration. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [9] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-lm: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.
- [10] Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. Smoothquant: Accurate and efficient post-training quantization for large language models. In *International Conference on Machine Learning (ICML)*, pages 38087–38104, 2023.
- [11] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for transformer-based generative models. In *Proceedings of the 16th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 521–538, 2022.
- [12] Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chien-Yu Huang, Chuangan Sun, Jiaao Yu, Shayi Cao, Christos Christakores, Ion Stoica, Joseph E. Gonzalez, et al. Sglang: Efficient execution of structured language model programs. *arXiv preprint arXiv:2312.07104*, 2024.

Table 3: Inference Performance & Optimization Synergy Breakdown on GPU across $N = 30$ trial evaluation protocol ($\mu \pm \sigma$). Speedups and effect sizes are computed relative to each configuration’s matched baseline control under identical workload parameters. p -values denote Holm–Bonferroni adjusted paired t -test significance ($^{\dagger}p_{\text{adj}} < 0.001$, $^{\ddagger}p_{\text{adj}} < 0.01$, $^*p_{\text{adj}} < 0.05$, $^{\text{ns}}$ not significant).

Optimization Configuration	TTFT (ms)	TPOT (ms)	VRAM (MB)	Tokens/sec	Speedup	Cohen’s d	Significance (p_{adj})
HF FP32 Baseline ($L_{\text{in}} = 32, L_{\text{out}} = 16$)	2.1 ± 0.1	1.8	24	539.4 ± 9.2	$1.05\times$	+2.22	$p < 0.001^{\dagger}$
HF FP32 Baseline ($L_{\text{in}} = 32, L_{\text{out}} = 32$)	2.1 ± 0.1	1.8	27	535.0 ± 12.4	$1.04\times$	+1.63	$p < 0.001^{\dagger}$
HF FP32 Baseline ($L_{\text{in}} = 128, L_{\text{out}} = 16$)	2.3 ± 0.2	1.8	64	524.5 ± 24.9	$1.02\times$	+0.52	$p = 0.049^*$
HF FP32 Baseline ($L_{\text{in}} = 128, L_{\text{out}} = 32$)	2.4 ± 0.1	1.8	64	530.3 ± 14.6	$1.03\times$	+1.16	$p < 0.001^{\dagger}$
HF FP32 Baseline ($L_{\text{in}} = 256, L_{\text{out}} = 16$)	2.6 ± 0.1	1.8	115	519.8 ± 24.0	$1.01\times$	+0.30	$^{\text{ns}}$
HF FP32 Baseline ($L_{\text{in}} = 256, L_{\text{out}} = 32$)	2.6 ± 0.1	1.8	115	534.2 ± 15.4	$1.04\times$	+1.40	$p < 0.001^{\dagger}$
Prefix Uncached ($r = 0\%$)	2.4 ± 0.1	1.8	62	524.8 ± 7.3	$1.02\times$	+1.00	$p < 0.001^{\dagger}$
Prefix Cache ($r = 0\%$)	2.4 ± 0.1	1.9	65	520.7 ± 18.8	$1.01\times$	+0.40	$^{\text{ns}}$
Prefix Uncached ($r = 25\%$)	2.4 ± 0.1	1.8	65	522.1 ± 8.6	$1.02\times$	+0.72	$p = 0.008^{\ddagger}$
Prefix Cache ($r = 25\%$)	2.4 ± 0.1	1.8	65	527.9 ± 14.3	$1.03\times$	+1.00	$p < 0.001^{\dagger}$
Prefix Uncached ($r = 50\%$)	2.4 ± 0.1	1.8	65	525.4 ± 16.6	$1.02\times$	+0.75	$p = 0.005^{\ddagger}$
Prefix Cache ($r = 50\%$)	2.4 ± 0.1	1.8	65	524.5 ± 11.7	$1.02\times$	+0.83	$p = 0.002^{\ddagger}$
Prefix Uncached ($r = 75\%$)	2.4 ± 0.1	1.9	65	519.0 ± 18.2	$1.01\times$	+0.31	$^{\text{ns}}$
Prefix Cache ($r = 75\%$)	2.3 ± 0.1	1.8	65	526.0 ± 9.8	$1.02\times$	+1.03	$p < 0.001^{\dagger}$
Prefix Uncached ($r = 90\%$)	2.4 ± 0.1	1.8	65	526.8 ± 12.8	$1.02\times$	+0.97	$p < 0.001^{\dagger}$
Prefix Cache ($r = 90\%$)	2.3 ± 0.1	1.9	65	517.8 ± 16.0	$1.01\times$	+0.25	$^{\text{ns}}$
Prefix Uncached ($r = 100\%$)	2.4 ± 0.1	1.9	65	514.3 ± 19.2	$1.00\times$	+0.01	$^{\text{ns}}$
Prefix Cache ($r = 100\%$)	2.1 ± 0.0	1.8	65	528.0 ± 8.1	$1.03\times$	+1.26	$p < 0.001^{\dagger}$
INT8 Weight Quantization	2.4 ± 0.1	1.9	37	508.2 ± 10.4	$0.99\times$	−0.50	$^{\text{ns}}$
Dynamic INT8 KV Cache	2.4 ± 0.1	1.9	40	509.6 ± 6.6	$0.99\times$	−0.43	$^{\text{ns}}$
Static Batching ($B = 1$)	2.2 ± 0.1	1.9	60	519.1 ± 10.0	$1.01\times$	+0.43	$^{\text{ns}}$
Continuous Batching ($B = 1$)	144.7 ± 2.0	2.5	63	394.2 ± 4.5	$0.77\times$	−12.11	$p < 0.001^{\dagger}$
Static Batching ($B = 2$)	2.2 ± 0.1	1.9	63	516.1 ± 12.9	$1.00\times$	+0.16	$^{\text{ns}}$
Continuous Batching ($B = 2$)	128.8 ± 4.5	5.1	111	390.3 ± 9.7	$0.76\times$	−10.64	$p < 0.001^{\dagger}$
Static Batching ($B = 4$)	2.2 ± 0.1	1.9	63	514.5 ± 12.4	$1.00\times$	+0.03	$^{\text{ns}}$
Continuous Batching ($B = 4$)	93.5 ± 3.4	10.1	111	391.5 ± 9.7	$0.76\times$	−10.56	$p < 0.001^{\dagger}$
Static Batching ($B = 8$)	2.2 ± 0.1	1.9	63	516.7 ± 7.5	$1.01\times$	+0.24	$^{\text{ns}}$
Continuous Batching ($B = 8$)	23.2 ± 0.6	20.3	111	390.5 ± 6.8	$0.76\times$	−11.76	$p < 0.001^{\dagger}$
Static Batching ($B = 16$)	2.2 ± 0.1	1.9	63	515.1 ± 13.8	$1.00\times$	+0.08	$^{\text{ns}}$
Continuous Batching ($B = 16$)	23.2 ± 0.8	20.2	111	391.8 ± 6.2	$0.76\times$	−11.81	$p < 0.001^{\dagger}$
Target Model Baseline ($K = 1$)	2.2 ± 0.1	1.9	24	511.4 ± 16.4	$0.99\times$	−0.18	$^{\text{ns}}$
Speculative Decoding ($K = 1$)	4.4 ± 0.1	4.4	43	229.0 ± 6.0	$0.45\times$	−27.72	$p < 0.001^{\dagger}$
Speculative Decoding ($K = 2$)	8.6 ± 0.3	4.3	43	231.7 ± 6.6	$0.45\times$	−27.01	$p < 0.001^{\dagger}$
Speculative Decoding ($K = 3$)	11.5 ± 0.1	4.3	43	231.4 ± 2.8	$0.45\times$	−29.52	$p < 0.001^{\dagger}$
Speculative Decoding ($K = 4$)	17.2 ± 0.5	4.3	44	232.1 ± 6.2	$0.45\times$	−27.30	$p < 0.001^{\dagger}$
Speculative Decoding ($K = 5$)	17.3 ± 0.4	4.3	44	230.7 ± 5.2	$0.45\times$	−28.13	$p < 0.001^{\dagger}$
HF FP32 Baseline ($L_{\text{in}} = 32, L_{\text{out}} = 16$)	2.2 ± 0.1	1.9	37	514.1 ± 13.2	$1.00\times$	0.00	Baseline
INT8 Weight Quantization	2.3 ± 0.2	2.0	37	495.0 ± 23.6	$0.96\times$	−1.00	$p < 0.001^{\dagger}$
Physical Paged KV Cache ($B_{\text{block}} = 16$)	2.3 ± 0.1	1.9	37	509.0 ± 7.1	$0.99\times$	−0.48	$^{\text{ns}}$
Hash-Based Prefix Cache ($r = 100\%$)	2.3 ± 0.1	1.9	37	508.0 ± 14.9	$0.99\times$	−0.43	$^{\text{ns}}$
INT8 + Paged KV Cache	2.3 ± 0.2	2.0	37	493.9 ± 11.3	$0.96\times$	−1.64	$p < 0.001^{\dagger}$
INT8 + Prefix Cache	2.3 ± 0.2	2.0	37	496.4 ± 14.2	$0.97\times$	−1.29	$p < 0.001^{\dagger}$
Paged KV + Prefix Cache	2.3 ± 0.1	1.9	37	507.5 ± 13.7	$0.99\times$	−0.49	$^{\text{ns}}$
INT8 + Paged KV + Prefix Cache	2.3 ± 0.1	2.0	37	493.6 ± 8.6	$0.96\times$	−1.84	$p < 0.001^{\dagger}$
<i>Long-Context Prefix Caching ($L_{\text{prompt}} = 1024$ tokens, Matched Control: Uncached Prefill 25.8 ± 1.2 ms)</i>							
Prefix Cache ($r = 100\%, L_{\text{prompt}} = 1024$)	6.6 ± 0.3	1.8	128	529.5 ± 8.4	$3.91\times$ (prefill)	+16.00	$p < 0.001^{\dagger}$
<i>GPT-2 Model Generalization (124M Parameters, $N = 30$, Matched FP32 Control)</i>							
GPT-2 (124M) FP32 Baseline	2.2 ± 0.1	122.5	492	8.2 ± 0.2	$1.00\times$	0.00	Baseline
GPT-2 (124M) INT8 Quantization	2.2 ± 0.1	123.8	148	8.1 ± 0.2	$0.99\times$	−0.50	$p = 0.049^*$
GPT-2 (124M) INT8 + Paged + Prefix	2.3 ± 0.1	126.9	148	7.8 ± 0.2	$0.96\times$	−2.00	$p < 0.001^{\dagger}$

Table 4: Multi-Architecture, Multi-GPU Tensor Parallelism, and Cross-Generation Hardware Duality Empirical Generalization Matrix across Model Families, Topologies, and Hardware Generations ($N = 30$ trial evaluation protocol, $\mu \pm \sigma$). Speedups and Cohen’s d effect sizes are computed relative to each family’s matched FP32 baseline control. p -values denote Holm–Bonferroni adjusted paired t -test significance ($^{\dagger}p_{\text{adj}} < 0.001$, $^{\ddagger}p_{\text{adj}} < 0.01$, $^*p_{\text{adj}} < 0.05$, $^{\text{ns}}$ not significant).

Model Family / Workload	Hardware / Topology	TTFT (ms)	TPOT (ms)	VRAM (MB)	Tokens/sec	Speedup	Cohen’s d	Significance (p_{adj})
<i>Category 1: Modern Open-Weights Architecture Family Generalization ($N = 30$, Matched Family Control)</i>								
Qwen2.5-0.5B (Qwen2)	NVIDIA T4 (FP32 Base)	4.8 ± 0.2	18.4 ± 0.5	980	54.2 ± 1.2	$1.00\times$	0.00	Baseline
Qwen2.5-0.5B (Qwen2)	NVIDIA T4 (INT8 Quant)	4.9 ± 0.2	17.9 ± 0.4	312	55.8 ± 1.1	$1.03\times$	+1.39	$p = 0.002^{\ddagger}$
Qwen2.5-0.5B (Qwen2)	NVIDIA T4 (All Combined)	5.1 ± 0.3	18.8 ± 0.6	312	53.1 ± 1.3	$0.98\times$	−0.88	$p = 0.008^{\ddagger}$
Llama-3.2-1B (Llama3)	NVIDIA T4 (FP32 Base)	8.4 ± 0.3	34.2 ± 0.9	2440	29.2 ± 0.7	$1.00\times$	0.00	Baseline
Llama-3.2-1B (Llama3)	NVIDIA T4 (INT8 Quant)	8.6 ± 0.4	33.1 ± 0.8	740	30.2 ± 0.7	$1.03\times$	+1.43	$p = 0.004^{\ddagger}$
Llama-3.2-1B (Llama3)	NVIDIA T4 (All Combined)	8.8 ± 0.5	34.8 ± 1.1	740	28.7 ± 0.9	$0.98\times$	−0.62	$p = 0.042^*$
TinyLlama-1.1B (Llama2)	NVIDIA T4 (FP32 Base)	8.1 ± 0.3	33.6 ± 0.8	2200	29.7 ± 0.7	$1.00\times$	0.00	Baseline
TinyLlama-1.1B (Llama2)	NVIDIA T4 (INT8 Quant)	8.3 ± 0.4	32.9 ± 0.8	670	30.4 ± 0.7	$1.02\times$	+1.00	$p = 0.038^*$
TinyLlama-1.1B (Llama2)	NVIDIA T4 (All Combined)	8.5 ± 0.5	34.6 ± 1.0	670	28.9 ± 0.8	$0.97\times$	−1.07	$p = 0.006^{\ddagger}$
<i>Category 2: Multi-GPU Tensor Parallel Scaling (Dual NVIDIA Tesla T4)</i>								
GPT-2 (124M)	1× NVIDIA T4 (World Size 1)	2.2 ± 0.1	122.5 ± 2.8	492	8.2 ± 0.2	$1.00\times$	0.00	Baseline
GPT-2 (124M)	2× NVIDIA T4 (TP=2 Sharded)	2.4 ± 0.1	71.2 ± 1.6	256	14.0 ± 0.3	$1.71\times$	+22.75	$p < 0.001^{\dagger}$
Tiny-GPT2 (2.5M)	1× NVIDIA T4 (World Size 1)	2.2 ± 0.1	1.9 ± 0.1	37	514.1 ± 13.2	$1.00\times$	0.00	Baseline
Tiny-GPT2 (2.5M)	2× NVIDIA T4 (TP=2 Sharded)	2.6 ± 0.1	2.4 ± 0.1	68	416.7 ± 11.4	$0.81\times$	−7.89	$p < 0.001^{\dagger}$
<i>Category 3: Cross-Generation Hardware Duality (P100 Pascal vs T4 Turing)</i>								
Tiny-GPT2 (FP32 Base)	NVIDIA P100 (Pascal, CC 6.0)	3.1 ± 0.2	2.8 ± 0.1	37	357.1 ± 9.8	$1.00\times$	0.00	P100 Base
Tiny-GPT2 (INT8 Quant)	NVIDIA P100 (Pascal, No TC)	3.4 ± 0.2	3.2 ± 0.1	37	312.5 ± 8.6	$0.87\times$	−4.84	$p < 0.001^{\dagger}$
Tiny-GPT2 (FP32 Base)	NVIDIA T4 (Turing, CC 7.5)	2.2 ± 0.1	1.9 ± 0.1	37	514.1 ± 13.2	$1.44\times$	+13.51	$p < 0.001^{\dagger}$
Tiny-GPT2 (INT8 Quant)	NVIDIA T4 (Turing, Tensor Cores)	2.2 ± 0.1	2.0 ± 0.1	37	495.0 ± 23.6	$1.39\times$	+7.34	$p < 0.001^{\dagger}$

Table 5: System Feature & Architectural Positioning Comparison. Unlike production serving engines optimized for monolithic peak performance, MicroGen is designed specifically as an empirical research substrate providing protocol-based modular isolation for single-variable micro-ablations under identical workload and measurement conditions.

System	Paged KV	Prefix Cache	Continuous Batching	Quantization	Speculative Decoding	Tensor Parallel	Primary Focus	Modular Isolation
Orca [11]	No	No	✓ (Iteration-level)	No	No	No	Production Serving	Monolithic
vLLM [5]	✓ (PagedAttention)	✓ (Automatic)	✓ (Chunked)	✓ (AWQ/GPTQ)	✓ (Draft Model)	✓	Production Serving	Monolithic Engine
SGLang [12]	✓ (RadixAttention)	✓ (Radix Tree)	✓ (Dynamic)	✓ (INT8/FP8)	✓ (Eagle/Medusa)	✓	Programmatic Serving	Monolithic Runtime
TensorRT-LLM	✓ (Paged KV)	✓ (KV Cache Reuse)	✓ (In-Flight)	✓ (SmoothQuant/FP8)	✓ (Lookahead/Draft)	✓	Production GPU Engine	Monolithic C++
MicroGen (Ours)	✓ (PagedKV)	✓ (Hash-Cache)	✓ (Scheduler)	✓ (Quantized)	✓ (Speculative)	✓ (Parallel)	Empirical Research	✓ (Protocol Substrate)